


MLP RL controller architectures used in the study

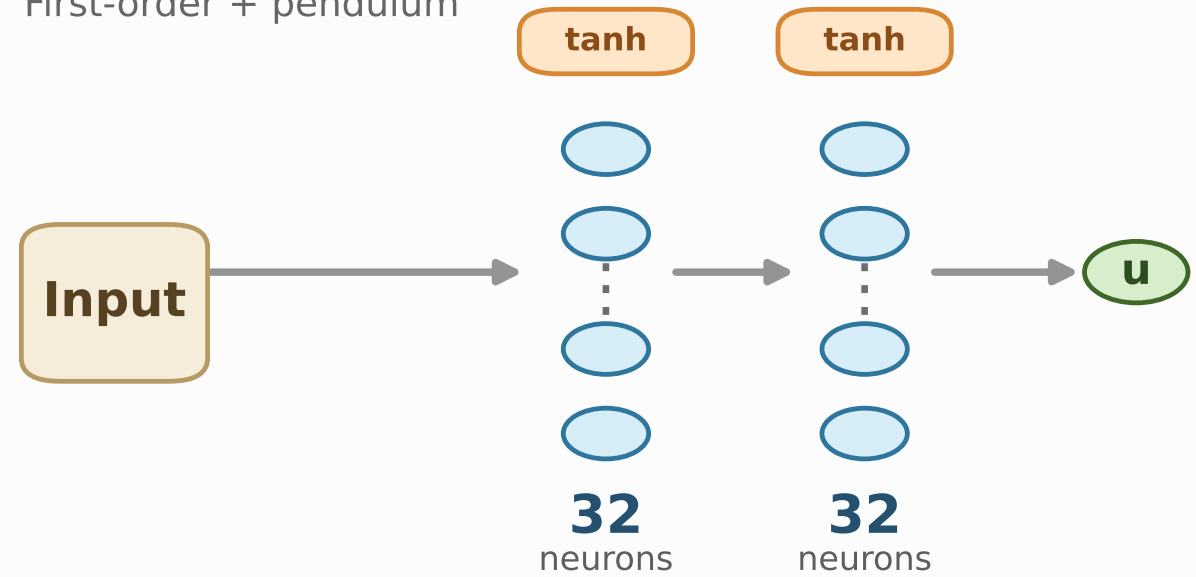
MLP 16-16

First-order + pendulum



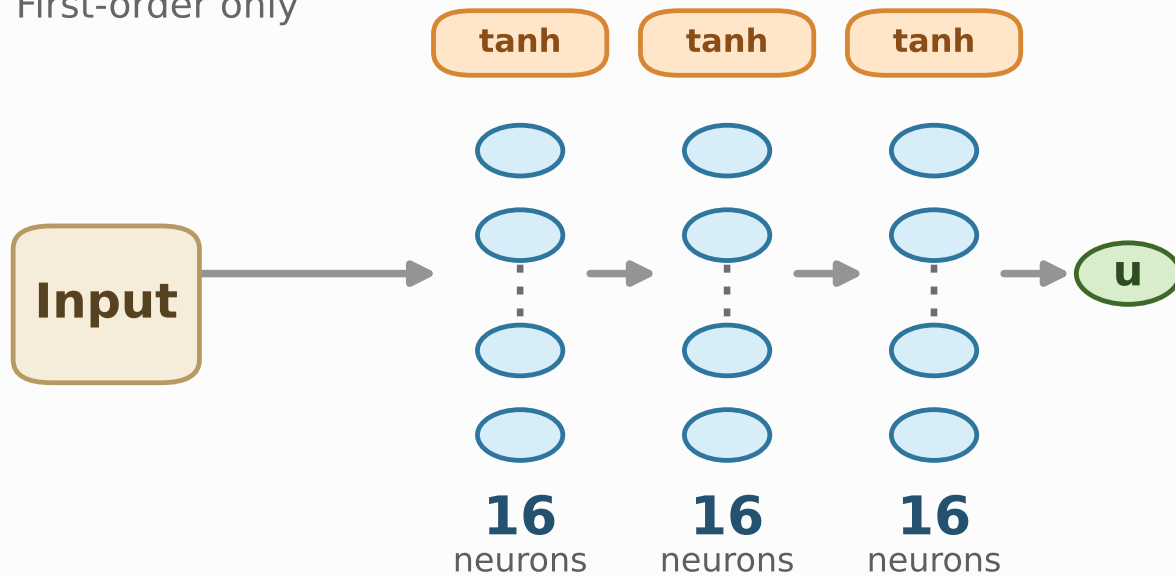
MLP 32-32

First-order + pendulum



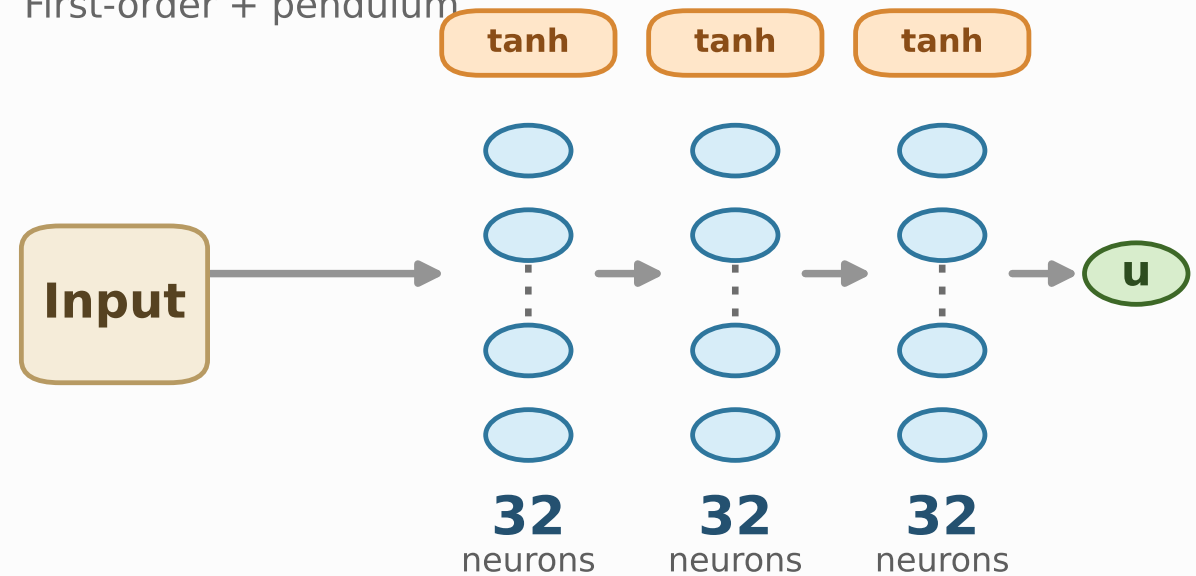
MLP 16-16-16

First-order only



MLP 32-32-32

First-order + pendulum



Hidden layers use tanh activation. In the main benchmark runs the final output is linear: $u = u_{\text{raw}}$.