


MLP RL controller architectures used in the study


MLP 16-16

First-order + pendulum




MLP 32-32

First-order + pendulum




MLP 16-16-16

First-order only



MLP 32-32-32

First-order + pendulum



Hidden layers use tanh activation. Final output is smoothly bounded via $u = \text{mid} + \text{half tanh}(u_{\text{raw}})$.